

```
import pandas as pd
import numpy as np
```

```
#Read files:
traindataset = pd.read_csv("Train.csv")
testdataset = pd.read_csv("Test.csv")
print (traindataset.head(10))
testdataset.head(10)
```

```

Item_Identifier  Item_Weight  ...  Outlet_Type  Item_Outlet_Sales
0      FDA15      9.300  ...  Supermarket Type1      3735.1380
1      DRC01      5.920  ...  Supermarket Type2      443.4228
2      FDN15     17.500  ...  Supermarket Type1     2097.2700
3      FDX07     19.200  ...  Grocery Store      732.3800
4      NCD19      8.930  ...  Supermarket Type1     994.7052
5      FDP36     10.395  ...  Supermarket Type2     556.6088
6      FDO10     13.650  ...  Supermarket Type1     343.5528
7      FDP10      NaN  ...  Supermarket Type3    4022.7636
8      FDH17     16.200  ...  Supermarket Type1    1076.5986
9      FDU28     19.200  ...  Supermarket Type1    4710.5350

```

[10 rows x 12 columns]

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_Outlet_Sales
0	FDW58	20.750	Low Fat	0.007565	Snack Foods	107.80
1	FDW14	8.300	reg	0.038428	Dairy	87.3
2	NCN55	14.600	Low Fat	0.099575	Others	241.7
3	FDQ58	7.315	Low Fat	0.015388	Snack Foods	155.0
4	FDY38	NaN	Regular	0.118599	Dairy	234.2
5	FDH56	9.800	Regular	0.063817	Fruits and Vegetables	117.1
6	FDL48	19.350	Regular	0.082602	Baking Goods	50.1
7	FDC48	NaN	Low Fat	0.015782	Baking Goods	81.0
8	FDN33	6.305	Regular	0.123365	Snack Foods	95.7
9	FDA36	5.985	Low Fat	0.005698	Baking Goods	186.8

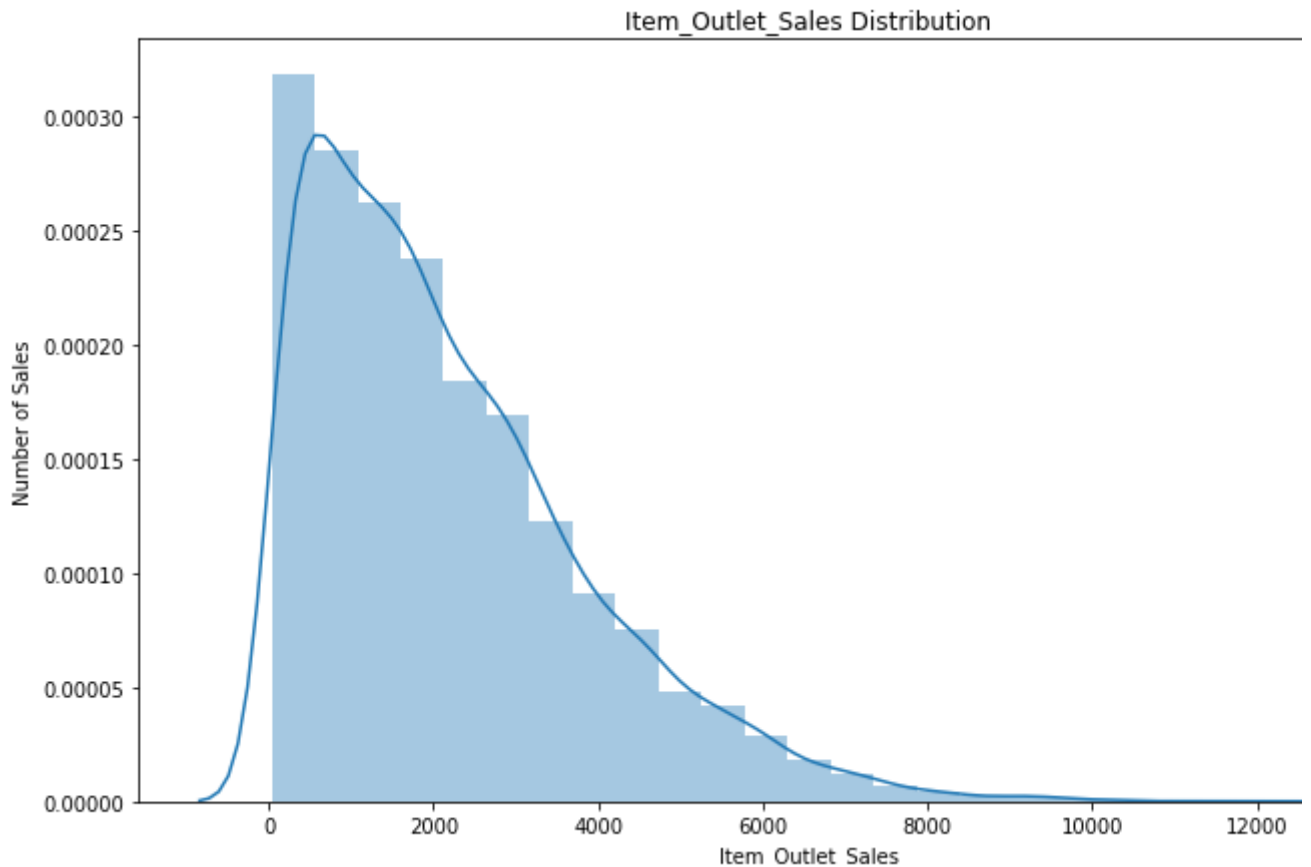
```
traindataset['source']='train'
testdataset['source']='test'
data = pd.concat([traindataset, testdataset],ignore_index=True)
print (traindataset.shape, testdataset.shape, data.shape)
```

```
print (traindataset.shape, testdataset.shape, data.shape)
```

```
(8523, 13) (5681, 12) (14204, 13)
```

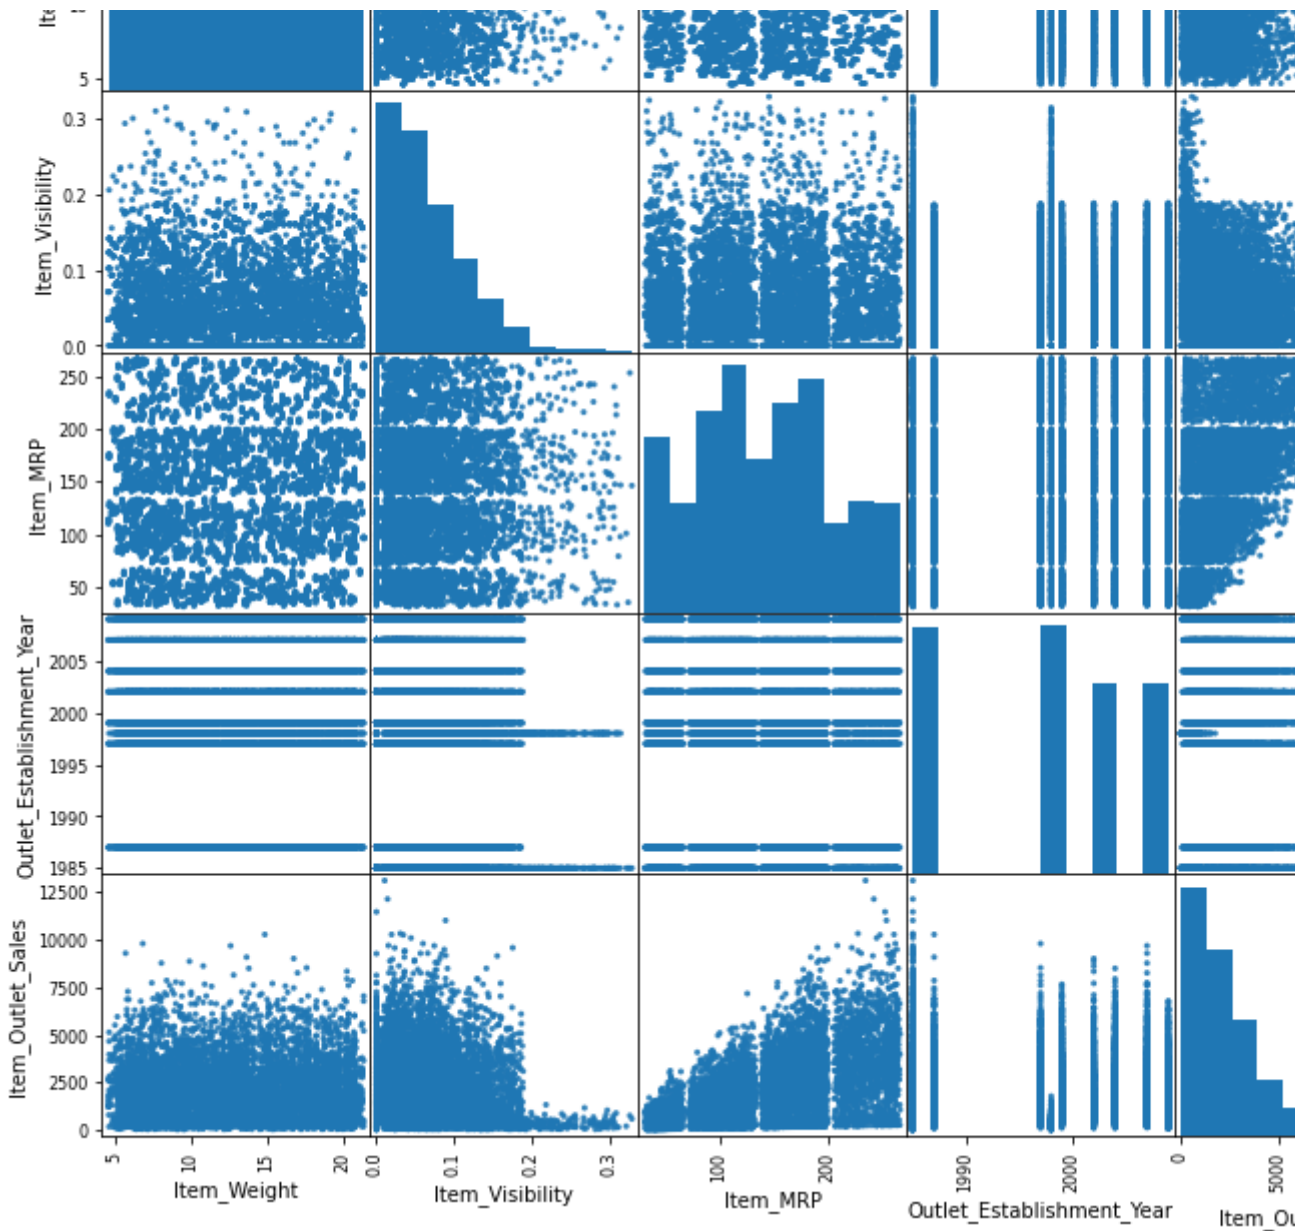
```
import matplotlib.pyplot as plt
import seaborn as sns
plt.figure(figsize=(12,7))
sns.distplot(traindataset.Item_Outlet_Sales, bins = 25)
plt.ticklabel_format(style='plain', axis='x', scilimits=(0,1))
plt.xlabel("Item_Outlet_Sales")
plt.ylabel("Number of Sales")
plt.title("Item_Outlet_Sales Distribution")
```

```
Text(0.5, 1.0, 'Item_Outlet_Sales Distribution')
```



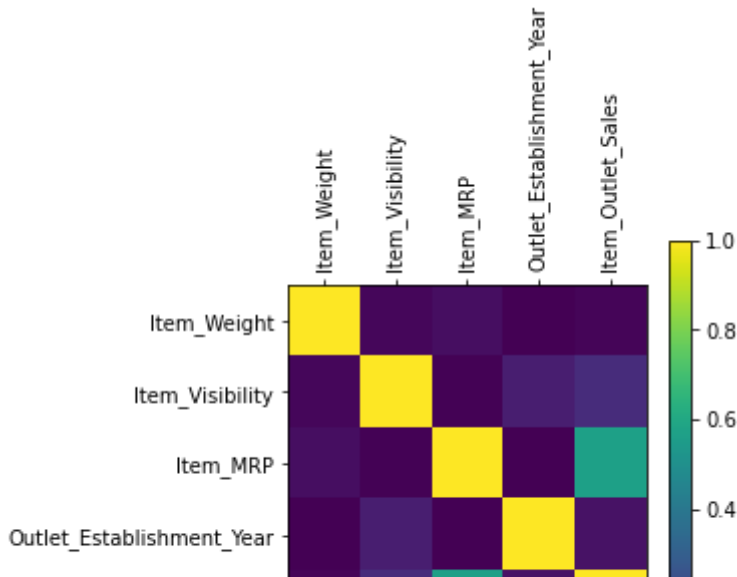
```
from pandas.plotting import scatter_matrix
scatter_matrix(data, alpha=0.9, figsize=(12,12), diagonal='hist')
plt.show()
pair_corr_coeff = data.corr()
print(pair_corr_coeff)
#pair_corr_coeff.abs().style.background_gradient()

plt.matshow(np.abs(pair_corr_coeff))
plt.colorbar()
plt.xticks(range(len(pair_corr_coeff.columns)), pair_corr_coeff.columns, rotation='vertical')
plt.yticks(range(len(pair_corr_coeff.columns)), pair_corr_coeff.columns)
plt.show()
```



	Item_Weight	...	Item_Outlet_Sales
Item_Weight	1.000000	...	0.014123
Item_Visibility	-0.015901	...	-0.128625
Item_MRP	0.036236	...	0.567574
Outlet_Establishment_Year	0.000645	...	-0.049135
Item_Outlet_Sales	0.014123	...	1.000000

[5 rows x 5 columns]





```
plt.figure(figsize=(12,7))
plt.xlabel("Item_Weight")
plt.ylabel("Item_Outlet_Sales")
plt.title("Item_Weight and Item_Outlet_Sales Analysis")
plt.plot(traindataset.Item_Weight, traindataset["Item_Outlet_Sales"],'.', alpha = 0.3)
```


☞ [`<matplotlib.lines.Line2D at 0x7f2c82330be0>`]



```
plt.figure(figsize=(12,7))
plt.xlabel("Item_Visibility")
plt.ylabel("Item_Outlet_Sales")
plt.title("Item_Visibility and Item_Outlet_Sales Analysis")
plt.plot(traindataset.Item_Visibility, traindataset["Item_Outlet_Sales"],".", alpha = 0.3)
```


☞

[<matplotlib.lines.Line2D at 0x7f2c824a6358>]



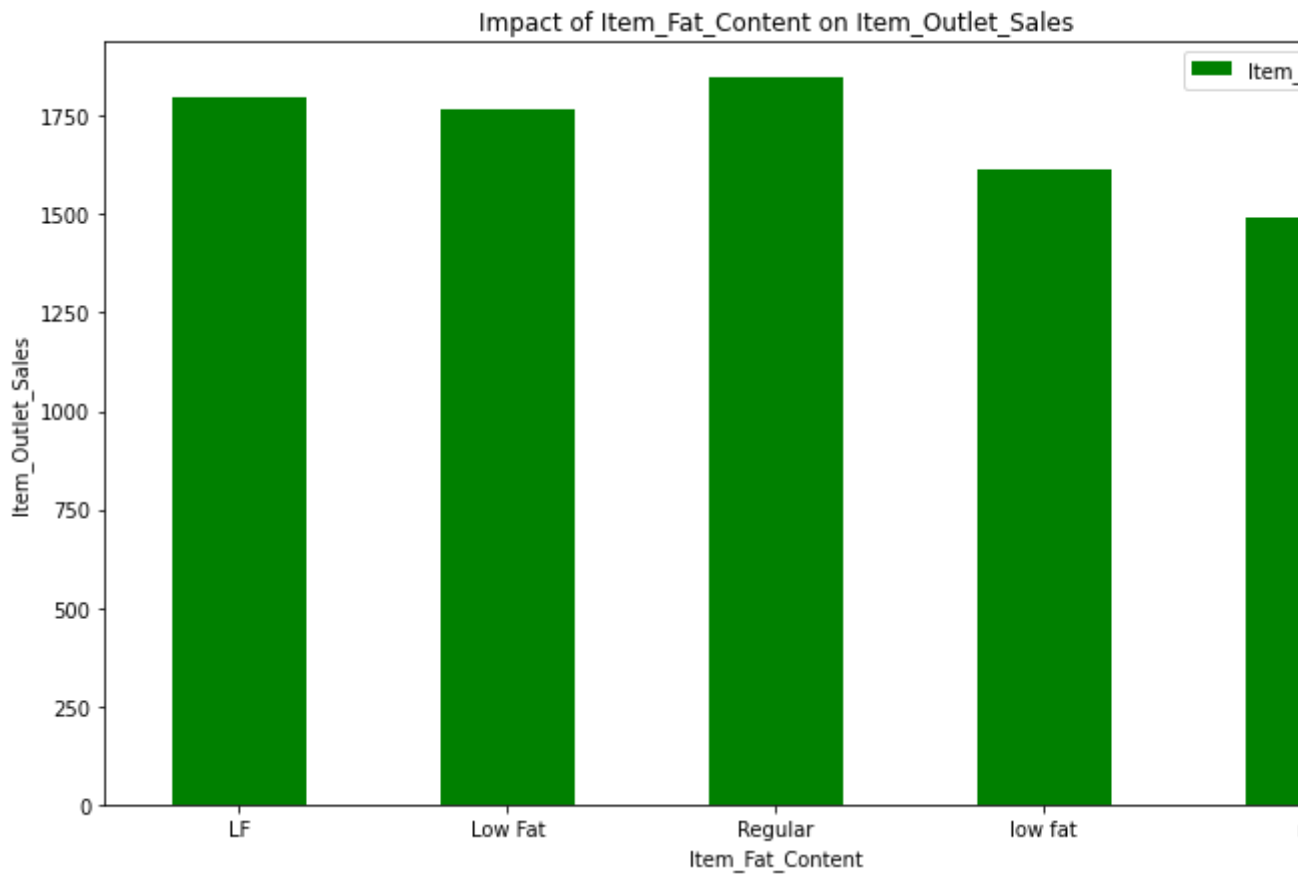
```
Outlet_Establishment_Year_pivot = traindataset.pivot_table(index='Outlet_Establishment_Year',
                                                             values='Item_Outlet_Sales',
                                                             aggfunc='sum')
Outlet_Establishment_Year_pivot.plot(kind='bar', color='red', figsize=(12,7))
plt.xlabel("Outlet_Establishment_Year")
plt.ylabel("Sum Item_Outlet_Sales")
plt.title("Impact of Outlet_Establishment_Year on Item_Outlet_Sales")
plt.xticks(rotation=0)
plt.show()
```





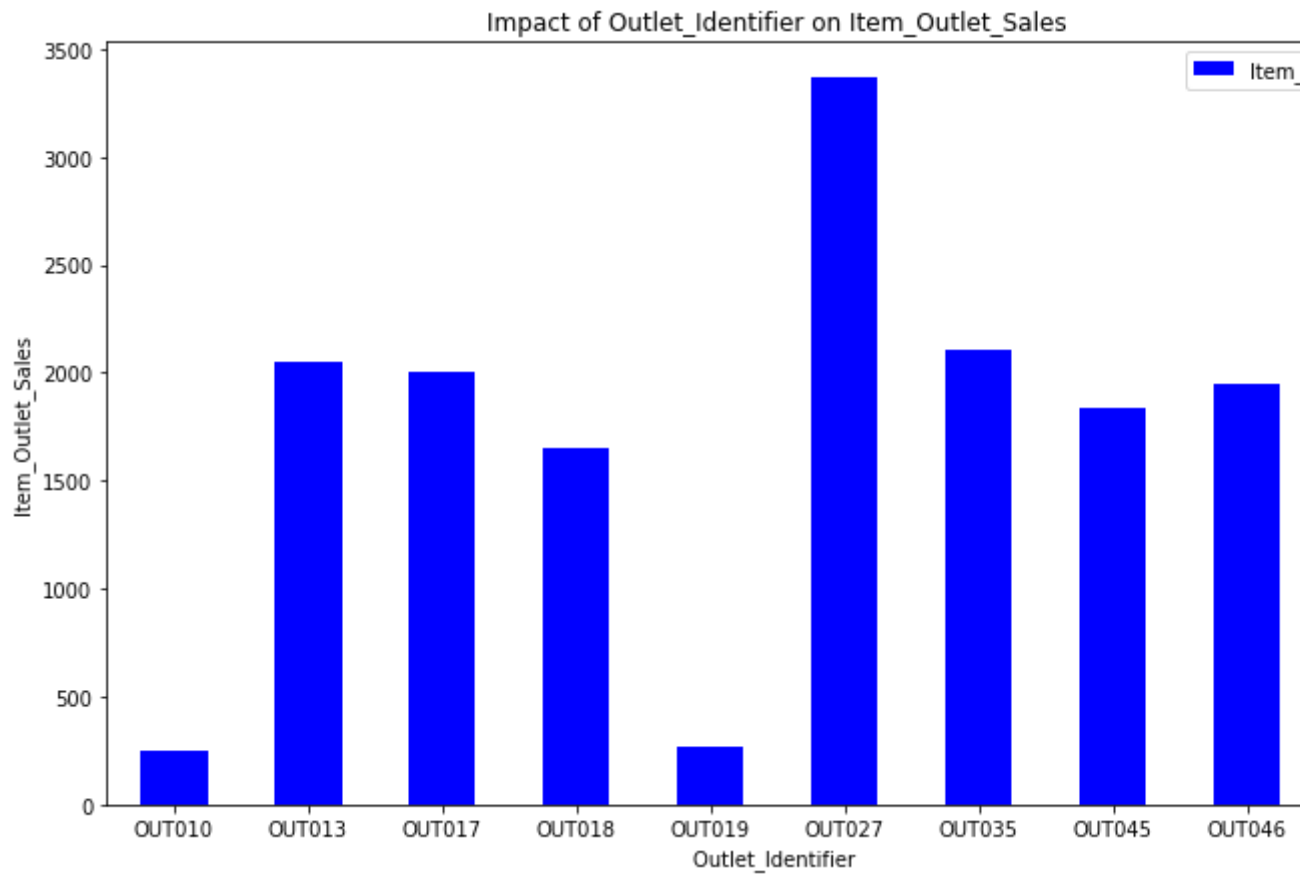
```
Item_Fat_Content_pivot =traindataset.pivot_table(index='Item_Fat_Content', values="Item_Ou
Item_Fat_Content_pivot.plot(kind='bar', color='green',figsize=(12,7))
plt.xlabel("Item_Fat_Content")
plt.ylabel("Item_Outlet_Sales")
plt.title("Impact of Item_Fat_Content on Item_Outlet_Sales")
plt.xticks(rotation=0)
plt.show()
```





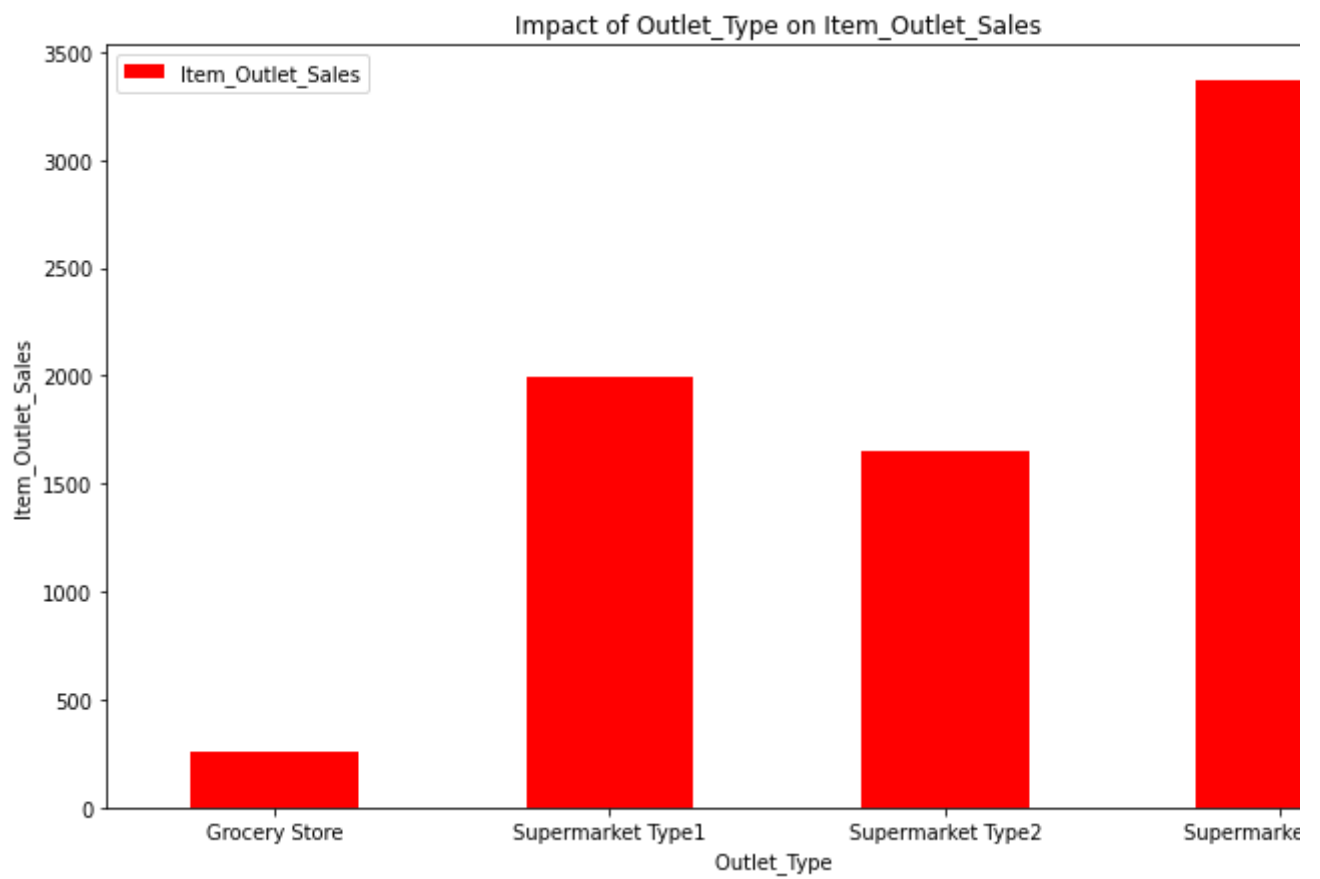
```
Outlet_Identifier_pivot = traindataset.pivot_table(index="Outlet_Identifier", values="Item_Outlet_Sales")
Outlet_Identifier_pivot.plot(kind="bar", color="blue", figsize=(12,7))
plt.xlabel("Outlet_Identifier ")
plt.ylabel("Item_Outlet_Sales")
plt.title("Impact of Outlet_Identifier on Item_Outlet_Sales")
plt.xticks(rotation=0)
plt.show()
```





```
Outlet_Type_pivot = traindataset.pivot_table(index='Outlet_Type', values="Item_Outlet_Sales")
Outlet_Type_pivot.plot(kind='bar', color='red',figsize=(12,7))
plt.xlabel("Outlet_Type ")
plt.ylabel("Item_Outlet_Sales")
plt.title("Impact of Outlet_Type on Item_Outlet_Sales")
plt.xticks(rotation=0)
plt.show()
```






```
print(data.apply(lambda x: sum(x.isnull())))  
sns.heatmap(data.isnull(),cbar=False,cmap='viridis',yticklabels=False)
```



```
Item_Identifier      0
Item_Weight         2439
Item_Fat_Content     0
Item_Visibility     0
Item_Type           0
Item_MRP            0
Outlet_Identifier    0
Outlet_Establishment_Year  0
Outlet_Size         4016
Outlet_Location_Type 0
Outlet_Type         0
Item_Outlet_Sales   5681
source              0
dtype: int64
<matplotlib.axes._subplots.AxesSubplot at 0x7f2c850d65c0>
```



```
data.describe()
```



	Item_Weight	Item_Visibility	Item_MRP	Outlet_Establishment_Year	Item_Outlet_Sales
count	11765.000000	14204.000000	14204.000000	14204.000000	14204.000000
mean	12.792854	0.065953	141.004977	1997.830681	12847.551815
std	4.652502	0.051459	62.086938	8.371664	2912.259036
min	4.555000	0.000000	31.290000	1985.000000	4.143000

```
data.apply(lambda x: len(x.unique()))
```

```

Item_Identifier      1559
Item_Weight          416
Item_Fat_Content      5
Item_Visibility     13006
Item_Type             16
Item_MRP             8052
Outlet_Identifier     10
Outlet_Establishment_Year  9
Outlet_Size           4
Outlet_Location_Type  3
Outlet_Type           4
Item_Outlet_Sales    3494
source                2
dtype: int64

```

```
#Filter categorical variables
```

```
categorical_columns = [x for x in data.dtypes.index if data.dtypes[x]=='object']
```

```
#Exclude ID cols and source:
```

```
categorical_columns = [x for x in categorical_columns if x not in ['Item_Identifier','Outlet_Identifier','source']]
```

```
#Print frequency of categories
```

```
for col in categorical_columns:
```

```
    print ('\nFrequency of Categories for variable %s'%col)
```

```
    print (data[col].value_counts())
```

```


```

Frequency of Categories for variable Item_Fat_Content

Low Fat	8485
Regular	4824
LF	522
reg	195
low fat	178

Name: Item_Fat_Content, dtype: int64

Frequency of Categories for variable Item_Type

Fruits and Vegetables	2013
Snack Foods	1989
Household	1548
Frozen Foods	1426
Dairy	1136
Baking Goods	1086
Canned	1084
Health and Hygiene	858
Meat	736
Soft Drinks	726
Breads	416
Hard Drinks	362
Others	280
Starchy Foods	269
Breakfast	186
Seafood	89

Name: Item_Type, dtype: int64

Frequency of Categories for variable Outlet_Size

Medium	4655
Small	3980
High	1553

Name: Outlet_Size, dtype: int64

Frequency of Categories for variable Outlet_Location_Type

Tier 3	5583
Tier 2	4641
Tier 1	3980

Name: Outlet_Location_Type, dtype: int64

Frequency of Categories for variable Outlet_Type

Supermarket Type1	9294
Grocery Store	1805
Supermarket Type3	1559
Supermarket Type2	1546

Name: Outlet_Type, dtype: int64

```

item_avg_weight = data.pivot_table(values= item_weight , index= item_identifier )
def impute_weight(cols):
    Weight = cols[0]
    Identifier = cols[1]

    if pd.isnull(Weight):
        return item_avg_weight['Item_Weight'][item_avg_weight.index == Identifier]
    else:
        return Weight
data['Item_Weight'] = data[['Item_Weight','Item_Identifier']].apply(impute_weight,axis=1).
print ('Final missing: %d'%sum(data['Item_Weight'].isnull()))

```

➞ Final missing: 0

```

#Import mode function:
from scipy.stats import mode

#Determining the mode for each
outlet_size_mode = data.pivot_table(values='Outlet_Size', columns='Outlet_Type',aggfunc=(lambda x: mode(x).mode[0]))
print ('Mode for each Outlet_Type:')
print (outlet_size_mode)

#Get a boolean variable specifying missing Item_Weight values
miss_bool = data['Outlet_Size'].isnull()

#Impute data and check #missing values before and after imputation to confirm

data.loc[miss_bool,'Outlet_Size'] = data.loc[miss_bool,'Outlet_Type'].apply(lambda x: outlet_size_mode[x])
print("final values that are missing are")
print (sum(data['Outlet_Size'].isnull()))

```

➞ Mode for each Outlet_Type:

Outlet_Type	Grocery Store	Supermarket Type1	Supermarket Type2	Supermarket Type3
Outlet_Size	Small	Small	Medium	Medium

final values that are missing are
0

```

sns.heatmap(data.isnull(),cbar=False,cmap='viridis',yticklabels=False)

```

➞

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f2c85147320>
```



```
data.pivot_table(values='Item_Outlet_Sales',index='Outlet_Type')
```



Item_Outlet_Sales	
Outlet_Type	
Grocery Store	339.828500
Supermarket Type1	2316.181148
Supermarket Type2	1995.498739
Supermarket Type3	3694.038558

```
visibility_item_avg = data.pivot_table(values='Item_Visibility',index='Item_Identifier')
def impute_visibility_mean(cols):
    visibility = cols[0]
    item = cols[1]

    if visibility == 0:
        return visibility_item_avg['Item_Visibility'][visibility_item_avg.index == item]
    else:
        return visibility

print ('Original #zeros: %d'%sum(data['Item_Visibility'] == 0))
data['Item_Visibility'] = data[['Item_Visibility','Item_Identifier']].apply(impute_visibility_mean,axis=1)
print ('Final #zeros: %d'%sum(data['Item_Visibility'] == 0))
```



```
Original #zeros: 879
Final #zeros: 0
```

```
data['Item_Visibility_MeanRatio'] = data.apply(lambda x: x['Item_Visibility']/visibility_item_avg[x['Item_Identifier']],axis=1)
data['Item_Visibility_MeanRatio'].describe()
```



```
count    14204.000000
mean         1.061884
std         0.235907
min         0.844563
25%         0.925131
50%         0.999070
75%         1.042007
max         3.010094
Name: Item_Visibility_MeanRatio, dtype: float64
```

```
data['Item_Type_Combined'] = data['Item_Identifier'].apply(lambda x: x[0:2])
#Rename them to more intuitive categories:
data['Item_Type_Combined'] = data['Item_Type_Combined'].map({'FD':'Food',
```

```
'NC':'Non-Consumable',
'DR':'Drinks'})
```

```
data['Item_Type_Combined'].value_counts()
```

```
Food          10201
Non-Consumable 2686
Drinks        1317
Name: Item_Type_Combined, dtype: int64
```

```
data['Outlet_Years'] = 2013 - data['Outlet_Establishment_Year']
data['Outlet_Years'].describe()
```

```
count    14204.000000
mean      15.169319
std       8.371664
min       4.000000
25%      9.000000
50%     14.000000
75%     26.000000
max     28.000000
Name: Outlet_Years, dtype: float64
```

```
print ('Original Categories:')
print (data['Item_Fat_Content'].value_counts())
```

```
print ('\nModified Categories:')
data['Item_Fat_Content'] = data['Item_Fat_Content'].replace({'LF':'Low Fat',
                                                             'reg':'Regular',
                                                             'low fat':'Low Fat'})
print (data['Item_Fat_Content'].value_counts())
```

```
Original Categories:
Low Fat    8485
Regular    4824
LF         522
reg        195
low fat    178
Name: Item_Fat_Content, dtype: int64
```

```
Modified Categories:
Low Fat    9185
Regular    5019
Name: Item_Fat_Content, dtype: int64
```

```
data.loc[data['Item_Type_Combined']=="Non-Consumable",'Item_Fat_Content'] = "Non-Edible"
data['Item_Fat_Content'].value_counts()
```

```
Low Fat    6499
Regular    5019
Non-Edible 2686
Name: Item_Fat_Content, dtype: int64
```

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
#New variable for outlet
```

```
data['Outlet'] = le.fit_transform(data['Outlet_Identifier'])
var_mod = ['Item_Fat_Content', 'Outlet_Location_Type', 'Outlet_Size', 'Item_Type_Combined', 'C
le = LabelEncoder()
for i in var_mod:
    data[i] = le.fit_transform(data[i])
```

```
s(data, columns=['Item_Fat_Content', 'Outlet_Location_Type', 'Outlet_Size', 'Outlet_Type', 'It
```

```
data.dtypes
```

```
Item_Identifier      object
Item_Weight          float64
Item_Visibility      float64
Item_Type            object
Item_MRP             float64
Outlet_Identifier     object
Outlet_Establishment_Year  int64
Item_Outlet_Sales    float64
source              object
Item_Visibility_MeanRatio float64
Outlet_Years         int64
Item_Fat_Content_0   uint8
Item_Fat_Content_1   uint8
Item_Fat_Content_2   uint8
Outlet_Location_Type_0 uint8
Outlet_Location_Type_1 uint8
Outlet_Location_Type_2 uint8
Outlet_Size_0        uint8
Outlet_Size_1        uint8
Outlet_Size_2        uint8
Outlet_Type_0        uint8
Outlet_Type_1        uint8
Outlet_Type_2        uint8
Outlet_Type_3        uint8
Item_Type_Combined_0  uint8
Item_Type_Combined_1  uint8
Item_Type_Combined_2  uint8
Outlet_0             uint8
Outlet_1             uint8
Outlet_2             uint8
Outlet_3             uint8
Outlet_4             uint8
Outlet_5             uint8
Outlet_6             uint8
Outlet_7             uint8
Outlet_8             uint8
Outlet_9             uint8
dtype: object
```

```
data[['Item_Fat_Content_0', 'Item_Fat_Content_1', 'Item_Fat_Content_2']].head(10)
```


	Item_Fat_Content_0	Item_Fat_Content_1	Item_Fat_Content_2
0	1	0	0
1	0	0	1
2	1	0	0
3	0	0	1
4	0	1	0
5	0	0	1
6	0	0	1
7	1	0	0
8	0	0	1
9	0	0	1

```
data.drop(['Item_Type','Outlet_Establishment_Year'],axis=1,inplace=True)
```

```
#Divide into test and train:
```

```
train = data.loc[data['source']=="train"]
```

```
test = data.loc[data['source']=="test"]
```

```
#Drop unnecessary columns:
```

```
test.drop(['Item_Outlet_Sales','source'],axis=1,inplace=True)
```

```
train.drop(['source'],axis=1,inplace=True)
```

```
#Export files as modified versions:
```

```
train.to_csv("train_modified.csv",index=False)
```

```
test.to_csv("test_modified.csv",index=False)
```

```
↳ /usr/local/lib/python3.6/dist-packages/pandas/core/frame.py:3997: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/stable/10min.html#copy-on-write>

```
mean_sales = train['Item_Outlet_Sales'].mean()
```

```
#Define a dataframe with TDs for submission:
```

```
""" Create a submission with test set predictions """
```

```
base1 = test[['Item_Identifier','Outlet_Identifier']]
base1['Item_Outlet_Sales'] = mean_sales
```

```
#Export submission file
base1.to_csv("alg0.csv",index=False)
```

↳ /usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:5: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/using_indexers.html

```
target = 'Item_Outlet_Sales'
IDcol = ['Item_Identifier','Outlet_Identifier']
from sklearn.model_selection import cross_validate
from sklearn.model_selection import cross_val_score
from sklearn import metrics
from sklearn.linear_model import LinearRegression, Ridge, Lasso
predictors = [x for x in train.columns if x not in [target]+IDcol]
# print predictors
alg1 = LinearRegression(normalize=True)

#Fit the algorithm on the data
alg1.fit(train[predictors], train[target])

#Predict training set:
train_predictions = alg1.predict(train[predictors])

#Perform cross-validation:
cv_score = cross_val_score(alg1, train[predictors], train[target], cv=20)
cv_score = np.sqrt(np.abs(cv_score))

#Print model report:
print ("\nModel Report")
print ("RMSE : %.4g" % np.sqrt(metrics.mean_squared_error(train[target].values, train_predictions)))
print ("CV Score : Mean - %.4g | Std - %.4g | Min - %.4g | Max - %.4g" % (np.mean(cv_score), np.std(cv_score), np.min(cv_score), np.max(cv_score)))

#Predict on testing data:
test[target] = alg1.predict(test[predictors])

#Export submission file:
IDcol.append(target)
submission = pd.DataFrame({ x: test[x] for x in IDcol})
submission.to_csv('alg1.csv', index=False)
coef1 = pd.Series(alg1.coef_, predictors).sort_values()
coef1.plot(title='Model Coefficients')
print("Accuracy is :")
```

```
print(alg1.score(train[predictors],train[target])*100)
```



Model Report

RMSE : 1127

CV Score : Mean - 0.7475 | Std - 0.0194 | Min - 0.6984 | Max - 0.7847


Accuracy is :

56.350509261468716

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:27: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/using_indexers.html



```
target = 'Item_Outlet_Sales'
IDcol = ['Item_Identifier','Outlet_Identifier']
from sklearn.model_selection import cross_validate
from sklearn.model_selection import cross_val_score
from sklearn import metrics
from sklearn.linear_model import LinearRegression, Ridge, Lasso
predictors = [x for x in train.columns if x not in [target]+IDcol]
alg2 = Ridge(alpha=0.05,normalize=True)

#Fit the algorithm on the data
alg2.fit(train[predictors], train[target])

#Predict training set:
train_predictions = alg2.predict(train[predictors])

#Perform cross-validation:
cv score = cross_val_score(alg2, train[predictors], train[target], cv=20)
```

```

cv_score = np.sqrt(np.abs(cv_score))

#Print model report:
print ("\nModel Report")
print ("RMSE : %.4g" % np.sqrt(metrics.mean_squared_error(train[target].values, train_predict)))
print ("CV Score : Mean - %.4g | Std - %.4g | Min - %.4g | Max - %.4g" % (np.mean(cv_score), np.std(cv_score), np.min(cv_score), np.max(cv_score)))

#Predict on testing data:
test[target] = alg2.predict(test[predictors])

#Export submission file:
IDcol.append(target)
submission = pd.DataFrame({ x: test[x] for x in IDcol})
submission.to_csv('alg2.csv', index=False)
coef2 = pd.Series(alg2.coef_, predictors).sort_values()
coef2.plot( title='Model Coefficients')
print("Accuracy is :")
print(alg2.score(train[predictors],train[target])*100)

```



Model Report

RMSE : 1129

CV Score : Mean - 0.7471 | Std - 0.01819 | Min - 0.7032 | Max - 0.7828


Accuracy is :

56.25404650408207

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:26: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/stable/10min.html>



```

target = 'Item_Outlet_Sales'
IDcol = ['Item_Identifier','Outlet_Identifier']
from sklearn.model_selection import cross_validate
from sklearn.model_selection import cross_val_score
from sklearn import metrics
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.tree import DecisionTreeRegressor

```

```

predictors = ['Item_MRP','Outlet_Type_0','Outlet_5','Outlet_Years']
alg2 = DecisionTreeRegressor(max_depth=8, min_samples_leaf=150)

```

```

alg3 = DecisionTreeRegressor(max_depth=8, min_samples_leaf=100)

#Fit the algorithm on the data
alg3.fit(train[predictors], train[target])

#Predict training set:
train_predictions = alg3.predict(train[predictors])

#Perform cross-validation:
cv_score = cross_val_score(alg3, train[predictors], train[target], cv=20)
cv_score = np.sqrt(np.abs(cv_score))

#Print model report:
print ("\nModel Report")
print ("RMSE : %.4g" % np.sqrt(metrics.mean_squared_error(train[target].values, train_predictions)))
print ("CV Score : Mean - %.4g | Std - %.4g | Min - %.4g | Max - %.4g" % (np.mean(cv_score), np.std(cv_score), np.min(cv_score), np.max(cv_score)))

#Predict on testing data:
test[target] = alg3.predict(test[predictors])

#Export submission file:
IDcol.append(target)
submission = pd.DataFrame({ x: test[x] for x in IDcol})
submission.to_csv('alg3.csv', index=False)
coef3 = pd.Series(alg3.feature_importances_, predictors).sort_values(ascending=False)
coef3.plot(title='Feature Importances')
print("Accuracy is : ")
print(alg3.score(train[predictors], train[target])*100)

```



```

Model Report
RMSE : 1071
CV Score : Mean - 0.7635 | Std - 0.02464 | Min - 0.7032 | Max - 0.8087
Accuracy is :
60.58974644993484
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:28: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/stable/10min.html#copy-on-write>

